

Experiment 8

Aim: To implement recommendation system on your dataset using the following machine learning techniques:

- Regression
- Classification
- Clustering
- Decision tree
- Anomaly detection
- Dimensionality Reduction
- Ensemble Methods

Theory:

- **Recommendation types:** Recommendation systems help users discover relevant items by predicting their preferences. The main types include content-based filtering, collaborative filtering, and hybrid methods. Content-based filtering recommends items similar to those the user has liked before, using item features such as category, description, or brand. Collaborative filtering, on the other hand, leverages the preferences of similar users or the patterns in user-item interactions without requiring item-specific data. It can be user-based or item-based, depending on whether it finds similarities between users or between items. Hybrid systems combine both methods to overcome limitations such as data sparsity or the cold-start problem, often resulting in better performance.
- **Recommendation measures:** Evaluating recommendation systems is essential to measure their effectiveness and accuracy. For prediction-based systems, RMSE (Root Mean Square Error) and MAE (Mean Absolute Error) are widely used to quantify the difference between predicted ratings and actual user ratings. Lower values indicate better accuracy. For ranking-based recommendations, metrics like precision, recall, and F1-score assess the system's ability to rank relevant items higher in a list. Additionally, Mean Reciprocal Rank (MRR) and Normalized Discounted Cumulative Gain (NDCG) are used to measure the quality of ranked recommendations. These metrics help determine how well the system delivers useful and personalized results to users.

Performance:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
from surprise import Dataset, Reader, SVD
from surprise.model_selection import train_test_split
from surprise import accuracy
```

```
orders = pd.read_csv("orders.csv")
order_products_prior = pd.read_csv("order_products__prior.csv")
order_products_train = pd.read_csv("order_products__train.csv")
products = pd.read_csv("products.csv")
aisles = pd.read_csv("aisles.csv")
departments = pd.read_csv("departments.csv")
```

```
products = products[["product_id", "aisle_id", "department_id", "product_name"]]
aisles = aisles[["aisle_id"]]
departments = departments[["department_id"]]
```

```
products = products.merge(aisles, on="aisle_id", how="left")
products = products.merge(departments, on="department_id", how="left")
```

```
order_products_prior = order_products_prior[["order_id", "product_id", "add_to_cart_order", "reordered"]]
order_products_train = order_products_train[["order_id", "product_id", "add_to_cart_order", "reordered"]]
```

```
order_products_prior = order_products_prior.merge(products, on="product_id", how="left")
order_products_train = order_products_train.merge(products, on="product_id", how="left")
```

```
order_products_prior = order_products_prior.sample(n=100000, random_state=42)
order_products_train = order_products_train.sample(n=50000, random_state=42)
```

```
all_orders = pd.concat([order_products_prior, order_products_train], ignore_index=True)
final_dataset = all_orders.merge(orders, on="order_id", how="left")
final_dataset.drop(columns=["eval_set", "aisle_id", "department_id"], inplace=True)
```

```
print("Final Dataset Columns:", final_dataset.columns)
print("Final Dataset Shape:", final_dataset.shape)
final_dataset.head()
```

```
Final Dataset Columns: Index(['order_id', 'product_id', 'add_to_cart_order', 'reordered',
                              'product_name', 'user_id', 'order_number', 'order_dow',
                              'order_hour_of_day', 'days_since_prior_order'],
                              dtype='object')
Final Dataset Shape: (150000, 10)
```

	order_id	product_id	add_to_cart_order	reordered	product_name	user_id	order_number	order_dow	order_hour_of_day	days_since_prior_order
0	3109255	34099	16	0	Crushed Red Chili Pepper	135284	9	0	19	8.0
1	301098	41950	5	0	Organic Tomato Cluster	7293	2	4	15	1.0
2	1181866	45066	8	0	Honeycrisp Apple	111385	2	1	17	8.0
3	1678630	8859	2	1	Natural Spring Water	147365	7	0	14	26.0
4	644090	24781	2	0	PODS Laundry Detergent, Ocean Mist Designed fo...	99290	7	0	19	30.0

```

from sklearn.impute import SimpleImputer

user_data = final_dataset.groupby("user_id").agg({
    "order_number": "mean",
    "days_since_prior_order": "mean",
    "add_to_cart_order": "mean"
}).reset_index()

imputer = SimpleImputer(strategy="mean")
user_data[["order_number", "days_since_prior_order", "add_to_cart_order"]] = imputer.fit_transform(
    user_data[["order_number", "days_since_prior_order", "add_to_cart_order"]]
)

scaler = StandardScaler()
scaled_features = scaler.fit_transform(user_data[["order_number", "days_since_prior_order", "add_to_cart_order"]])

kmeans = KMeans(n_clusters=5, random_state=42, n_init=10)
user_data["Cluster"] = kmeans.fit_predict(scaled_features)

print("Cluster Counts:\n", user_data["Cluster"].value_counts())

```

```

Cluster Counts:
Cluster
2      30293
1      22321
4      16973
0      11447
3       6634
Name: count, dtype: int64

```

This code clusters users based on their shopping behavior to uncover distinct patterns among customer segments

Collaborative Filtering Using Matrix Factorization (SVD):

```

cf_data = final_dataset[["user_id", "product_id", "reordered"]].dropna()

reader = Reader(rating_scale=(0, 1))
data = Dataset.load_from_df(cf_data, reader)
trainset, testset = train_test_split(data, test_size=0.2, random_state=42)

svd = SVD(n_factors=50, random_state=42)
svd.fit(trainset)

predictions = svd.test(testset)
rmse = accuracy.rmse(predictions)
print("Collaborative Filtering RMSE:", rmse)

```

```

RMSE: 0.4800
Collaborative Filtering RMSE: 0.47996185938330604

```

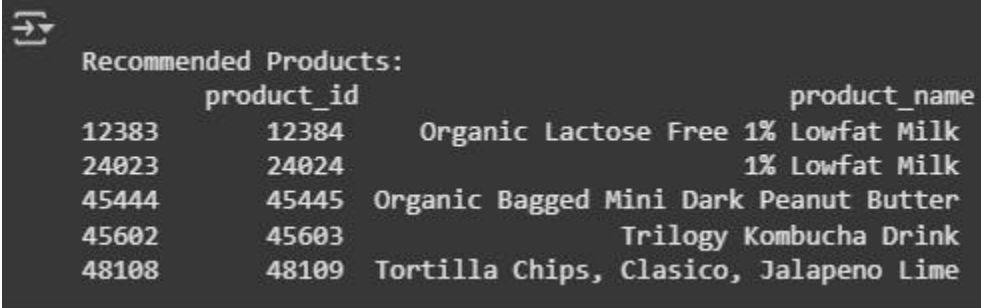
This code implements collaborative filtering to recommend products based on user-product interactions. RMSE (Root Mean Square Error) to measure the model's accuracy in predicting reorder behavior.

The model achieved a Collaborative Filtering RMSE of 0.47996.

Generating Personalized Product Recommendations:

```
def recommend_products(user_id, num_recommendations=5):
    unique_products = final_dataset["product_id"].unique()
    predicted_ratings = [(product, svd.predict(user_id, product).est) for product in unique_products]
    top_products = sorted(predicted_ratings, key=lambda x: x[1], reverse=True)[:num_recommendations]
    recommended_product_ids = [prod[0] for prod in top_products]
    recommended_products = products[products["product_id"].isin(recommended_product_ids)][["product_id", "product_name"]]
    return recommended_products

recommendations = recommend_products(user_id=1)
print("\nRecommended Products:\n", recommendations)
```



Recommended Products:

	product_id	product_name
12383	12384	Organic Lactose Free 1% Lowfat Milk
24023	24024	1% Lowfat Milk
45444	45445	Organic Bagged Mini Dark Peanut Butter
45602	45603	Trilogy Kombucha Drink
48108	48109	Tortilla Chips, Clasico, Jalapeno Lime

The above function leverages the trained SVD collaborative filtering model to recommend products to a specific user. When executed for `user_id=1`, the model recommended five specific items based on the user's shopping behavior: Organic Lactose Free 1% Lowfat Milk, 1% Lowfat Milk, Organic Bagged Mini Dark Peanut Butter, Trilogy Kombucha Drink, and Tortilla Chips, Clasico, Jalapeno Lime.

Conclusion:

In this experiment, we explored how to build a recommendation system by applying machine learning techniques, with a particular focus on clustering. We used K-Means to group users based on behavioral traits like how often they order, how recently they've ordered, and their typical cart activity. This helped us identify distinct user segments with shared shopping patterns. To push it further, we brought in collaborative filtering using Matrix Factorization (SVD), which gave solid predictive results, with a low RMSE of 0.47996, showing it can effectively predict what users are likely to reorder. The recommendation engine wasn't just accurate, it actually suggested relevant products like Organic Lactose Free 1% Lowfat Milk and Trilogy Kombucha, proving it could deliver useful, personalized results. By combining historical purchase behavior with clustering, we were able to better understand user needs and make smarter recommendations. Overall, the hybrid approach of segmentation plus collaborative filtering showed that blending techniques leads to more accurate and user-focused recommendation systems.

